

Lauren Hyoseo Yoon

laurenhyoon@caltech.edu | [LinkedIn](#) | PhD, Computation & Neural Systems, Caltech

RESEARCH

I engineer **efficient, stateful foundation models**, attacking two bottlenecks that keep today's models fragmented and forgetful:

- **Pillar I — Alignment** (*space*): force independently-formed representations into one shared, low-dimensional space.
- **Pillar II — Active inference-time learning** (*time*): let a system direct its own test-time compute, choosing what to explore and what to keep so that compute compounds.

EDUCATION

Ph.D., Computation & Neural Systems (CNS) | California Institute of Technology 2024 – present

Advisor: Yisong Yue.

Bachelor's, Mathematics & Neuroscience (Honors) | Northwestern University 2020 – 2024

High School Diploma | Korean Minjok Leadership Academy

Selective gifted boarding school, South Korea | International Philosophy Olympiad, National Representative (South Korea).

SELECTED RESEARCH

Pillar I · Multimodal Alignment forcing independently-formed representations into one shared space

Escaping Plato's Cave: JAM for Aligning Independently Trained Vision and Language Models ([link](#))

PMLR / NeurIPS 2025 (UniReps)

Aligns *frozen*, independently trained vision and language backbones with small per-modality autoencoders that meet at a shared bottleneck. A novel **Spread Loss** turns the Platonic Representation Hypothesis from something you *detect* into a training *objective* — matching a fine-tuned CLIP at $\sim 1/3$ of the FLOPs. *Google DeepMind; co-advised by Been Kim & Yisong Yue.*

Synthesizing Artificial Perception with a Wireless Transcranial Optogenetic Array ([link](#))

Nature Neuroscience 2025

Writes spatiotemporally *patterned* percepts into cortex via an addressable micro-LED array and **decodes them to behavioral choice**; discrimination scales with stimulation-site distance (PCA-denoised). Installing a representation and reading out what the animal made of it.

Striatal Ensemble Activity in an Innate Naturalistic Behavior ([link](#))

eLife 2023

A streaming **perception-to-representation pipeline**: lifts markerless video to 3D pose, then aligns it with unsupervised **population-level** neural trajectories (NMF) — movement and neural dynamics in one shared low-dimensional space. The same alignment problem as the AI work, in biology.

Wireless Multilateral Optofluidic Microsystems for Programmable Optogenetics & Photopharmacology ([link](#))

Nature Communications 2022

An **online connector system** for the brain — a GUI front-end bound to a high-frequency biological-signal backend over an NFC/RF link, driving up to 256 battery-free implants in real time, with behavior read back from video. The systems-engineering instinct that carries forward.

Pillar II · Active Inference-Time Learning directing test-time compute — what to explore, what to keep

Knowledge-Centric Self-Improvement (KCSI) ([link](#))

COLM 2026 ([CBW](#))

Inverts agent-centric self-improvement: agents stay generic and **disposable**, and the only thing that persists is a shared **curated knowledge base**. Transient agents attempt tasks and debate across them; what survives is distilled into reusable bundles. Higher solve rate at substantially lower cost (ARC, Polyglot, SWE-Bench Pro), 47.2% on Terminal-Bench, and bundles that transfer **zero-shot** across model families. *Current direction.*

VendiEvolve: Towards Diversity-Aware LLM-Guided Evolutionary Search ([link](#))

ICML 2026 (Human-AI Co-Creativity)

LLM-guided evolutionary search collapses onto a single program family, losing the stepping-stones to better solutions. VendiEvolve

puts the **Vendi Score** — the effective number of distinct programs — into the loop as a control signal steering selection toward diversity. **Diversity as a system control signal**; the population is the product, not just the single winner.

When Models Choose the Question: Pedagogical Constraints in Bottom-Up Multi-Agent Inquiry (link)

Preprint · arXiv 2026

Hand LLM agents a philosophical narrative and the freedom to ask their own questions — no facilitator, no aggregator — and inquiry quietly collapses back toward the framing the text already supplied. Freedom to *ask* is not freedom to *think otherwise*: the **ECO framework** names the ceiling on self-directed exploration when a system chooses what to ask.

EXPERIENCE

Student Researcher | Google DeepMind / Brain Jun 2024 – May 2025

Multimodal alignment of frozen vision–language models (JAM).

Research Intern, Biomedical Computing (Health Futures) | Microsoft Research Jun – Sep 2023

Foundation models of human DNA; proposed a genotype-imputation fine-tuning task and an efficient raw-data handling framework.

Research Scientist | Querrey Simpson Institute for Bioelectronics, Northwestern Fall 2023 – Jun 2024

End-to-end analysis pipelines for artificial-perception experiments (optogenetic/neural signals).

Research Scientist | Kozorovitskiy Lab, Northwestern Sep 2021 – Jun 2024

Senior thesis: interpretable generative models of behavior from electrophysiology + video; co-advised by Ann Kennedy (Feinberg).

Research Scientist | Computational Neuroscience Lab, Feinberg School of Medicine (Northwestern) Mar 2022 – Jan 2024

Python/ML behavior-tracking pipeline (MARS) characterizing circuit dysfunction in a Parkinson’s mouse model.

Research Assistant | Social Identity & Morality Lab, New York University Jan – Sep 2021

Built an automated social-media (Twitter API) program to test the backfire effect in fake-news discernment.